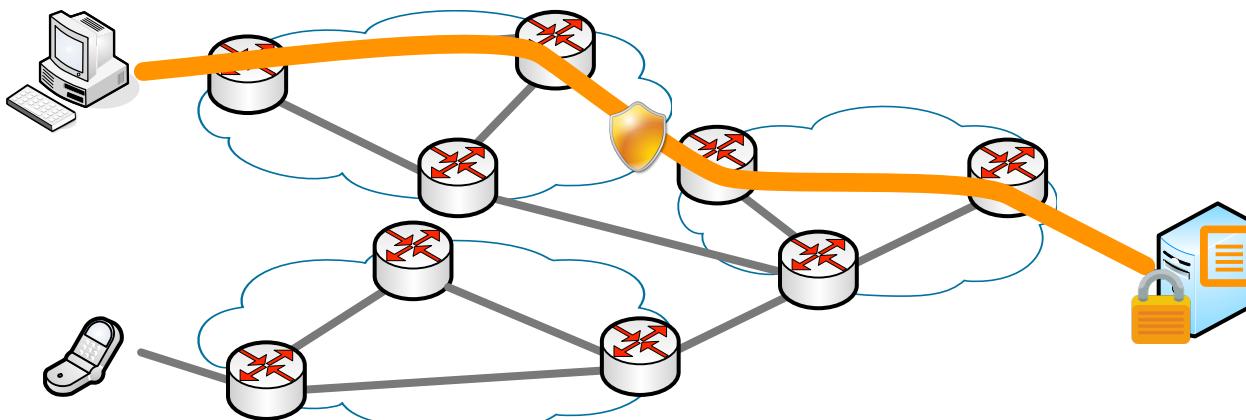


# Usable Security of Named Data Networking

Yingdi Yu

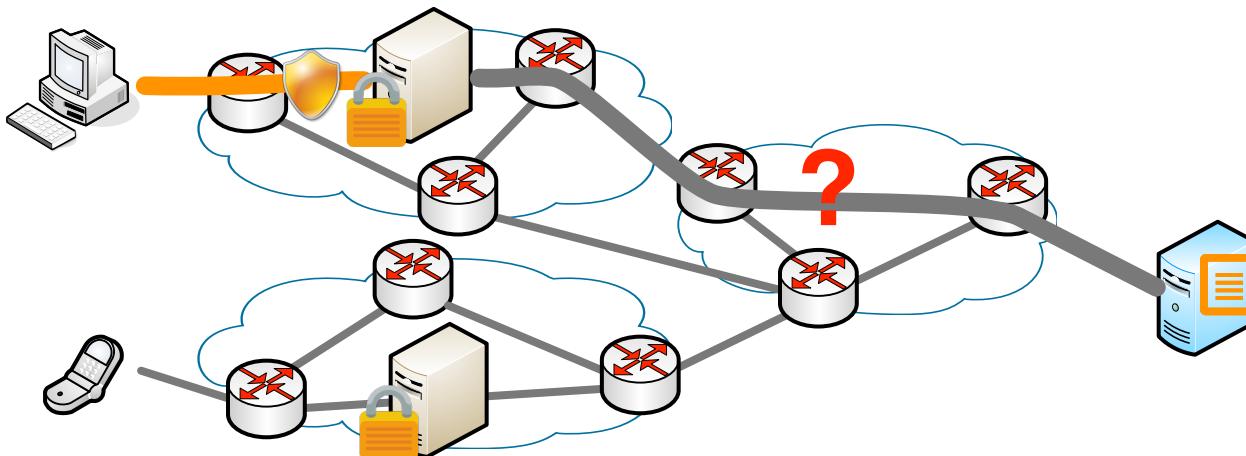
# Traditional communication model of Internet

- Speaking to a host
  - end-to-end channel
- Communication security
  - container-based authenticity: X.509, Certificate Authority
  - channel-based confidentiality: IPSec, TLS/SSL



# New communication vs. Old security

- Content Distribution Network (CDN)
  - multiple containers to secure
  - no end-to-end channel



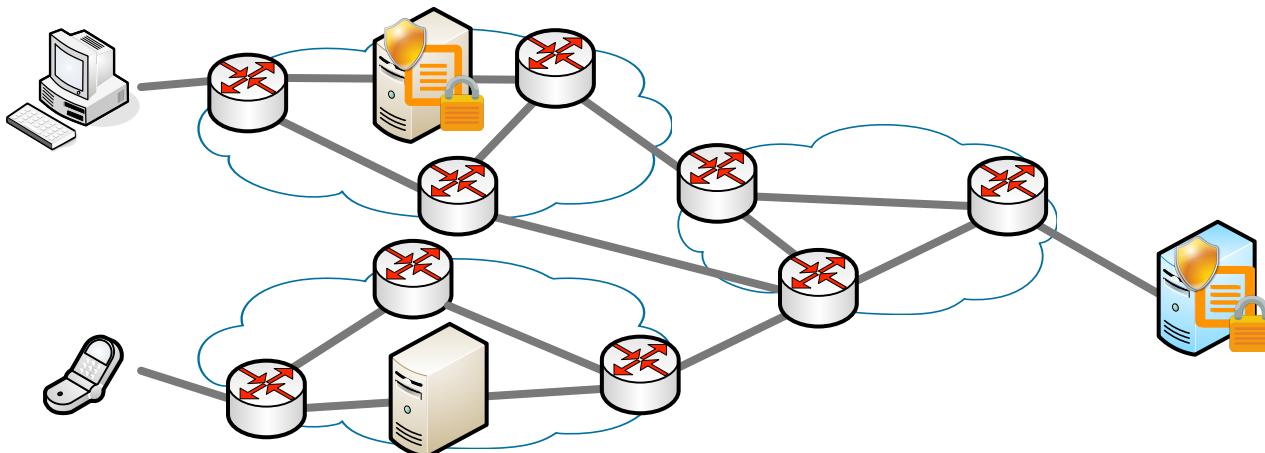
# New communication vs. Old security

- Delay Tolerant Network (DTN)
  - temporary data container
  - no instantaneous end-to-end channel



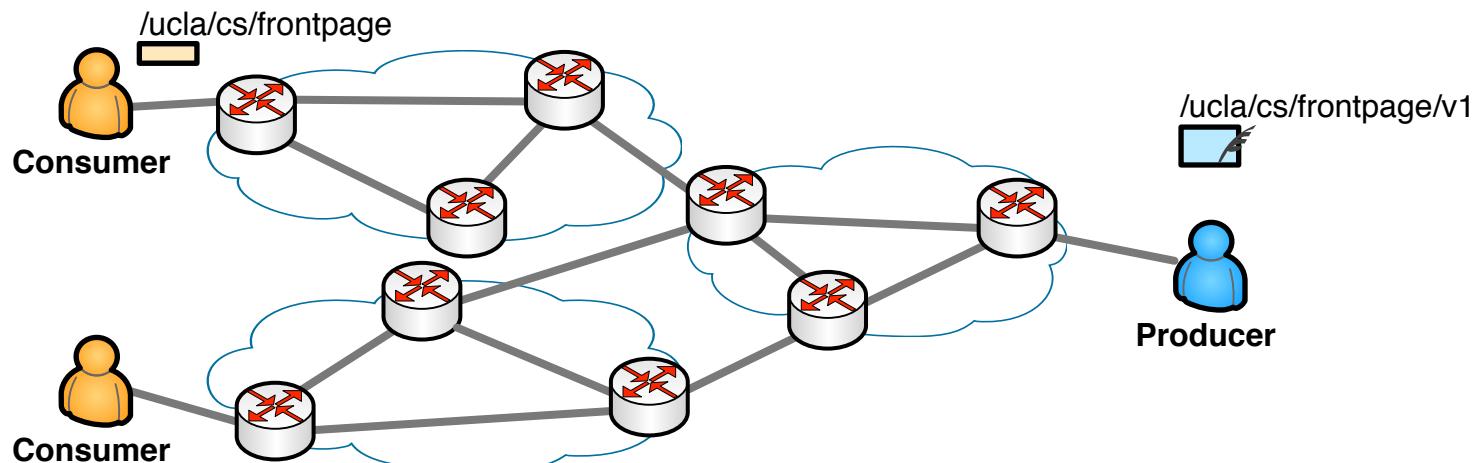
# New security model is desired!

- No trustworthy container, no end-to-end encrypted channel
- Data-centric security: let's secure **data** directly!
  - authenticate data rather than container
  - encrypt data instead of channel



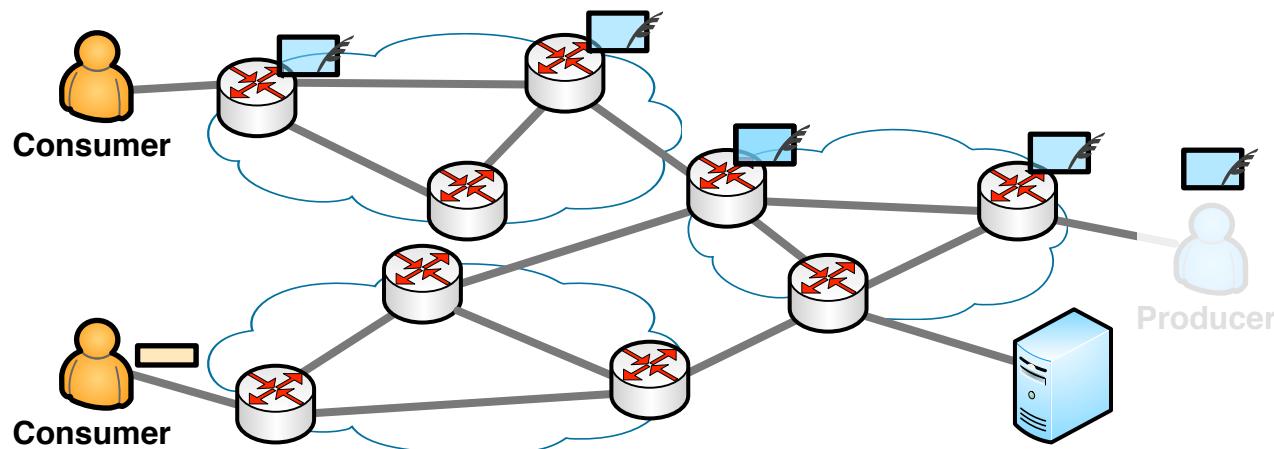
# Named Data Networking

- Data-centric communication primitives
  - retrieve data by name rather than container address
  - **Interest Packet**: expressed by consumer, forwarded according to name
  - **Data Packet**: made by producer, forwarded along reverse path



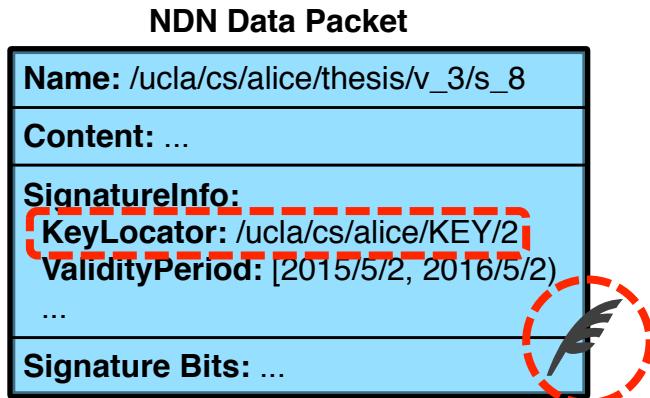
# Efficient & flexible data delivery

- Data can be picked anywhere
  - in-network caching
- Does not require instantaneous communication
  - producer can go offline
  - store pre-created data in third party storage



# Built-in data authenticity

- Per packet signature
  - privilege separation: different data signed by different keys
- Retrieve public key as data
  - same authentication procedure
- Data carrying public key is a certificate
  - more powerful



| NDN Certificate  |
|--|
| Name: /ucla/cs/alice/KEY/2   |
| Content:<br>6d:32:8d:23:a9:b0:89:...   |
| SignatureInfo:<br>SignatureType: RSA-SHA256<br>KeyLocator: /ucla/cs/KEY/7<br>ValidityPeriod: [2015/1/1, 2017/1/1)<br>... |
| Signature Bits:<br>cd:ca:70:72:7b:ff:a8:...  |

| X 509 Certificate                  |
|------------------------------------|
| Subject Name                       |
| Subject Public Key Info            |
| Certificate<br>Signature Algorithm |
| Issuer Name                        |
| Validity Period                    |
| Certificate Signature              |

# But how to utilize those features?

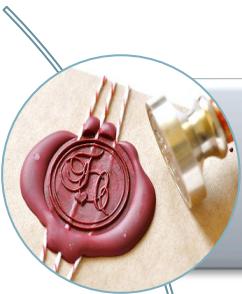
- Developers turn off security as the first step
  - fake signature
  - skip authentication
  - wish no one is eavesdropping



- Can we make security easier for developers?
  - automate data authentication
  - automate data encryption
  - minimize maintenance overhead



# Outline



Automating Data-Centric Authenticity

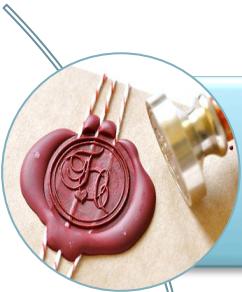


Authenticating Long-Lived Data

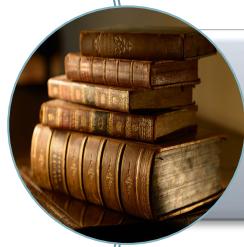


Automating Data-Centric Confidentiality

# Outline



**Automating Data-Centric Authenticity**



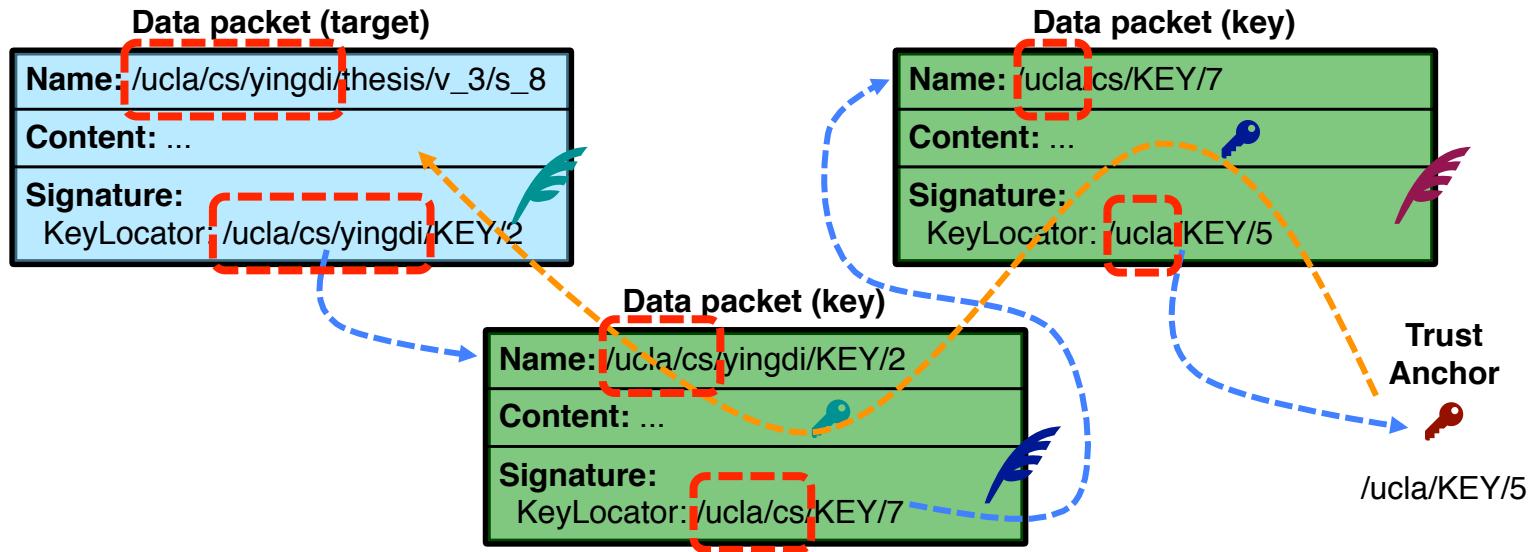
**Authenticating Long-Lived Data**



**Automating Data-Centric Confidentiality**

# Trust chain

- Recursively retrieve key until reach a **trust anchor**
  - a pre-trusted key
- Constrained by **trust derivation rules**
  - is data (or key) signed by a trusted producer (or issuer) ?
- Validate signature



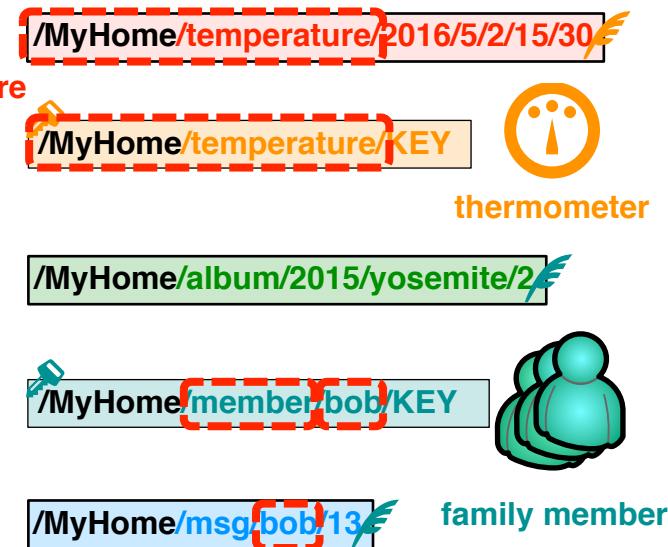
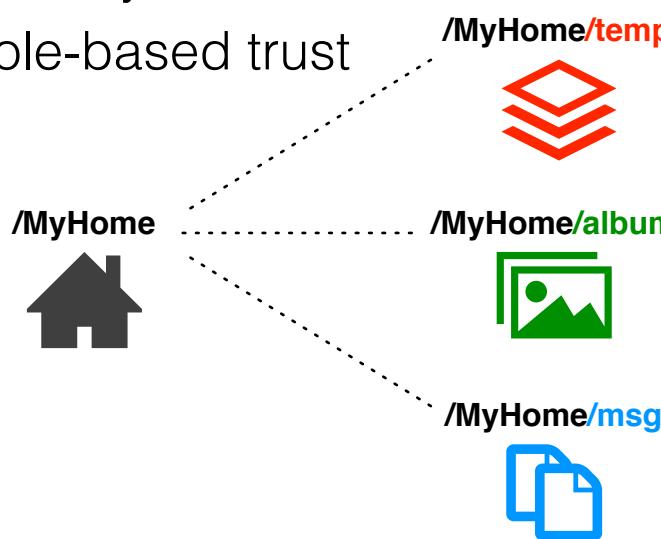
# Diversity of trust models

- Trust model could be simple in some cases



- Application specific in general

- capability-based trust
- identity-based trust
- role-based trust



# NDN insight

- Name is a general expression
  - can refer to identity, capability, role, ...
- Any trust model can be expressed as a list of relationships between data name and key name
- Data authentication easily if we have
  - a name-based policy language to express trust model
  - a library to perform authentication according to the policy

**Schematize the trust**

**Automate data authentication**

# Describe trust relationship in name

- Relationship between data and key names

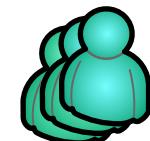
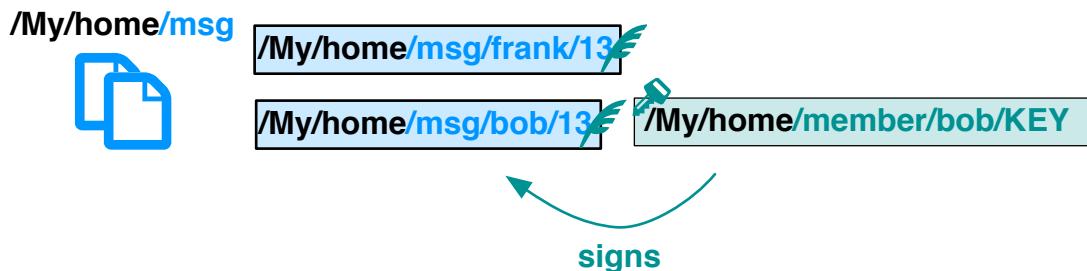
/My/home/msg/bob/13 ← /My/home/member/bob/KEY  
/My/home/msg/alice/15 ← /My/home/member/alice/KEY

- Generalized as name pattern

home\_prefix + "home" + "msg" + user + msg\_id  
home\_prefix + "home" + "member" + user + "KEY" → signs

- Regex-based syntax

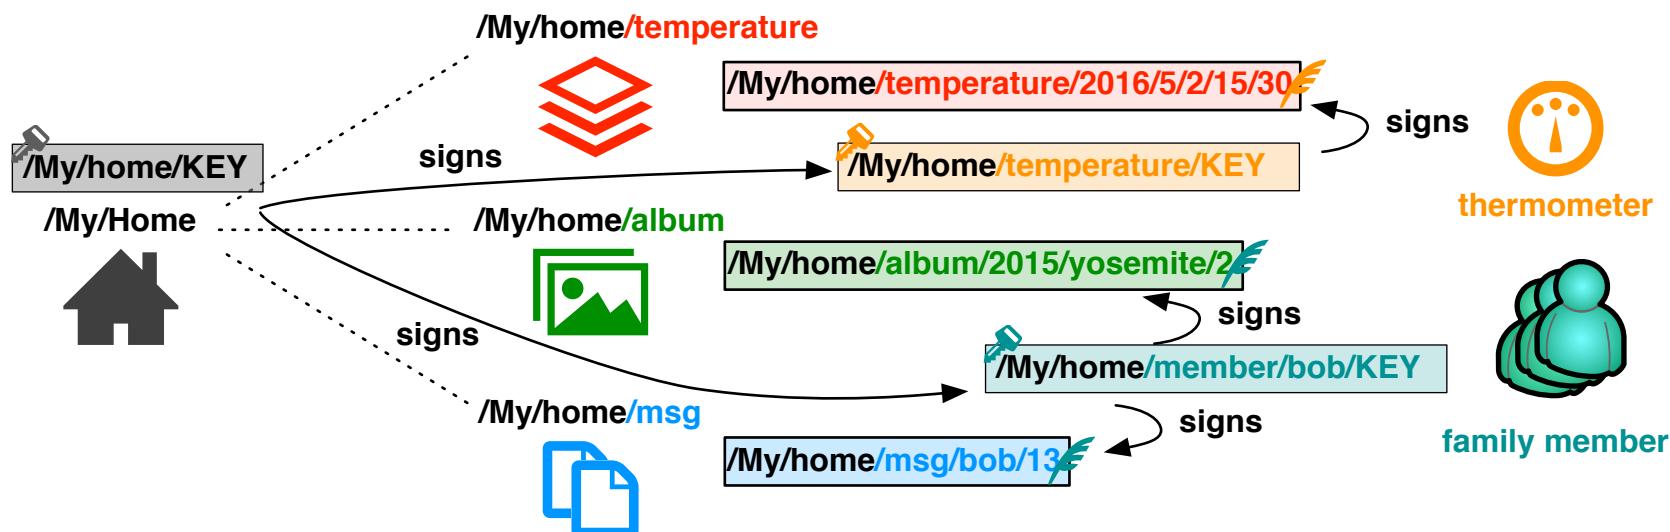
(<>\*)<home><msg>([user\_id])<> → signs  
\1<home><member>\2<KEY>



family member

# Trust schema

| Rule ID | Data Name                                     | Key Name                        |
|---------|---|---------------------------------|
| msg     | ( $\diamond^*$ )<home><msg>([user])<>         | \1<home><member>\2<KEY>         |
| album   | ( $\diamond^*$ )<home><album><><><>           | \1<home><member>[user]<KEY>     |
| temp    | ( $\diamond^*$ )<home><temperature><><><><><> | \1<home><temperature><KEY>      |
| member  | ( $\diamond^*$ )<home><member>([user])<KEY>   | \1<home><KEY>                   |
| therm   | ( $\diamond^*$ )<home><temperature><KEY>      | \1<home><KEY>                   |
| root    | ( $\diamond^*$ )<home><KEY>                   | /My/home/KEY 30:b4:82:9c:45:... |



# Trust chain construction

| Rule ID | Data Name  | Key Name                        |
|---------|--|---------------------------------|
| msg     | ( $\Delta$ * <home> <msg> ([user]) <>            | \1 <home> <member> \2 <KEY>     |
| album   | ( $\Delta$ * <home> <album> <> <> <>             | \1 <home> <member> [user] <KEY> |
| temp    | ( $\Delta$ * <home> <temperature> <> <> <> <> <> | \1 <home> <temperature> <KEY>   |
| member  | ( $\Delta$ * <home> <member> ([user]) <KEY>      | \1 <home> <KEY>                 |
| therm   | ( $\Delta$ * <home> <temperature> <KEY>          | \1 <home> <KEY>                 |
| root    | ( $\Delta$ * <home> <KEY>                        | /My/home/KEY 30:b4:82:9c:45:... |

**Data packet (target)**

|   |
|---|
| Name: /My/home/album/2015/yosemite/2              |
| Content: ...                                      |
| Signature:<br>KeyLocator: /My/home/member/bob/KEY |

**Data packet (key)**

|   |
|---|
| Name: /My/home/member/bob/KEY                     |
| Content: ...                                      |
| Signature:<br>KeyLocator: /My/home/member/bob/KEY |

**Data packet (key)**

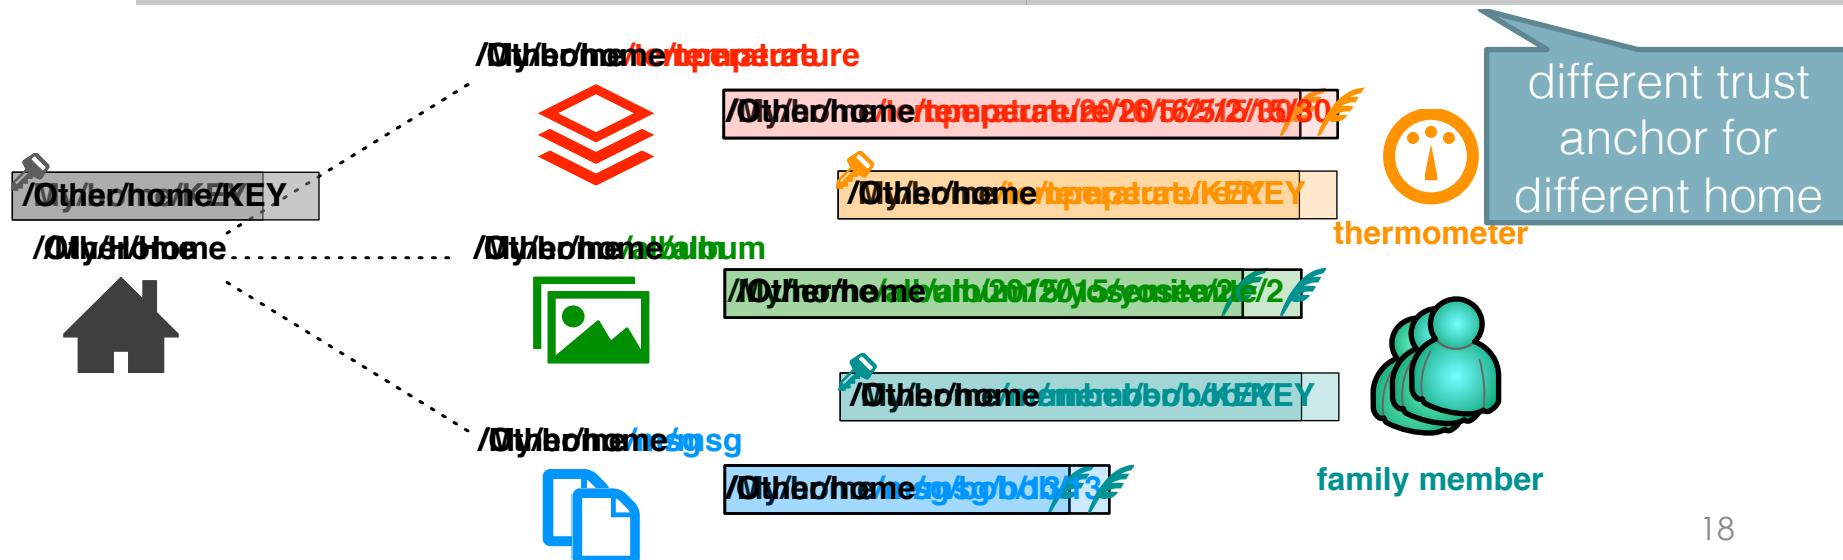
|  |
|--|
| Name: /My/home/member/bob/KEY          |
| Content: ...                           |
| Signature:<br>KeyLocator: /My/home/KEY |

Trust  
Anchor  


/My/home/KEY

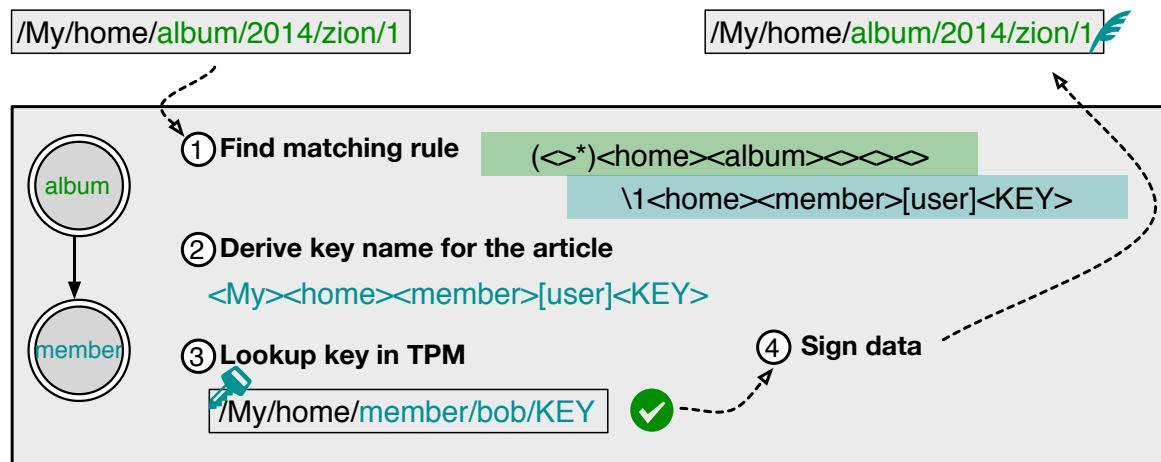
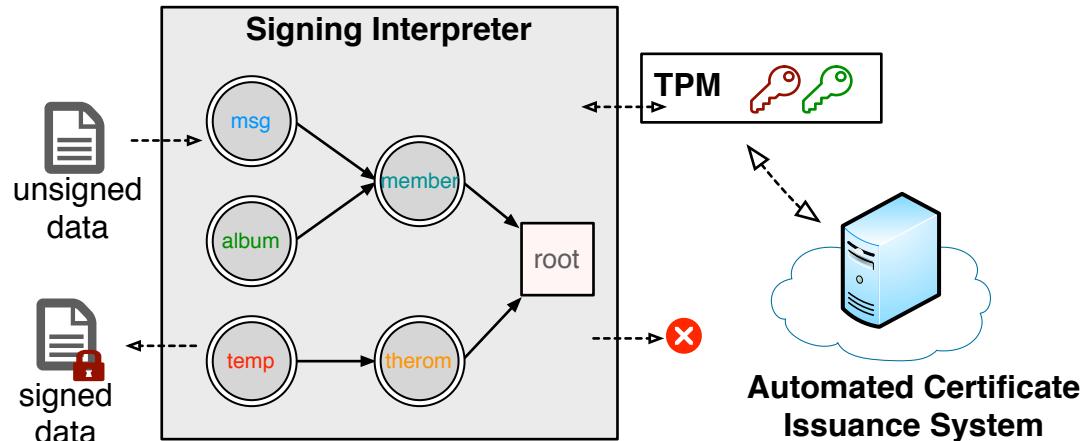
# Re-usability

| Rule ID | Data Name                                     | Key Name                           |
|---------|---|------------------------------------|
| msg     | ( $\diamond^*$ )<home><msg>([user])<>         | \1<home><member>\2<KEY>            |
| album   | ( $\diamond^*$ )<home><album><><><>           | \1<home><member>[user]<KEY>        |
| temp    | ( $\diamond^*$ )<home><temperature><><><><><> | \1<home><temperature><KEY>         |
| member  | ( $\diamond^*$ )<home><member>([user])<KEY>   | \1<home><KEY>                      |
| therm   | ( $\diamond^*$ )<home><temperature><KEY>      | \1<home><KEY>                      |
| root    | ( $\diamond^*$ )<home><KEY>                   | /Other/home/KEY 9c:45:30:b4:82:... |



# Automated Signing

- Signing Interpreter
- Determine signing key
- Request certificate if needed



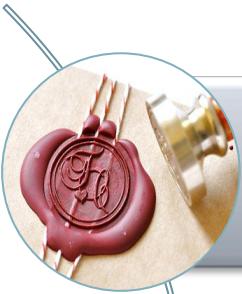
# Implementation

- Available in all the NDN platform libraries
  - ndn-cxx: <http://www.github.com/named-data/ndn-cxx>
  - NDN-CCL: <http://named-data.net/codebase/platform/ndn-ccl/>
- Powers data and interest authentication in:
  - NFD: NDN Forwarding
  - NLSR: NDN Link State Routing Protocol
  - NDNS: NDN Domain Name System
  - Repo-ng: NDN Data Repository
  - ChronoChat: server-less multi-party chat application over NDN

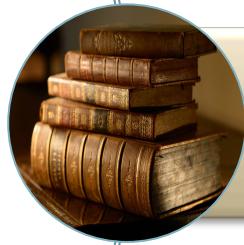
# Summary

- Trust schema is a **general expression** of NDN trust model
  - can be executed by any entity
- Trust schema is in **text format**
  - can be distributed as data packets
- A trust schema represents a **security design pattern**
  - regulate the behavior of applications
  - a set of common trust schemas

# Outline



Automating Data-Centric Authenticity



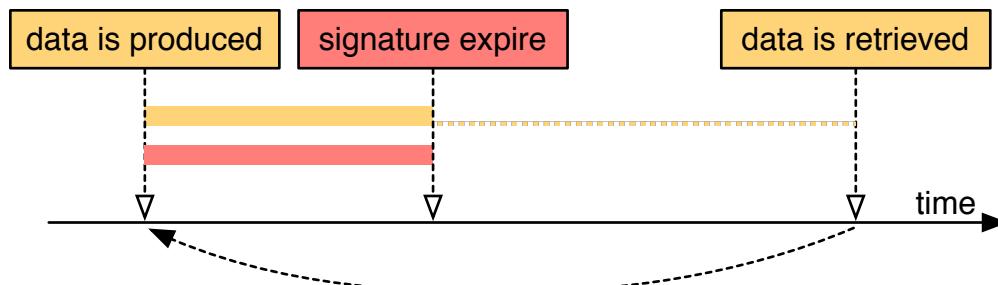
Authenticating Long-Lived Data



Automating Data-Centric Confidentiality

# Lifetime Mismatch

- Data lifetime is usually longer than its signature
  - crypto algorithm, key compromise, ...
- Periodical re-signing is not the solution
  - will not scale in long term
  - data may outlive its producer
  - not a problem in channel based security
- After fact validation
  - verify signature validity at the time of production



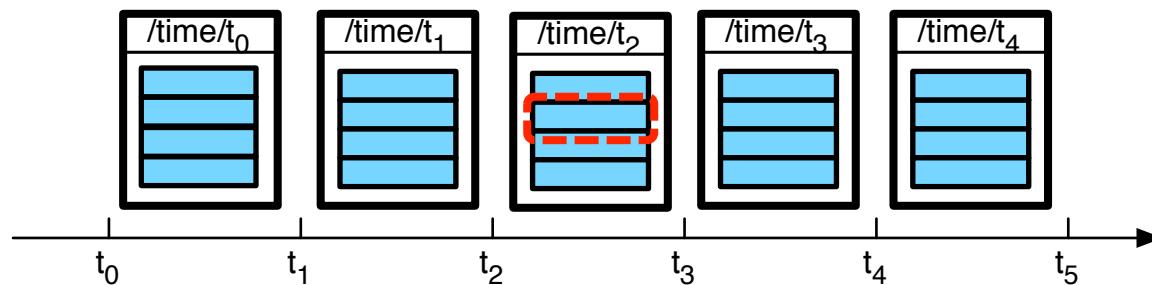
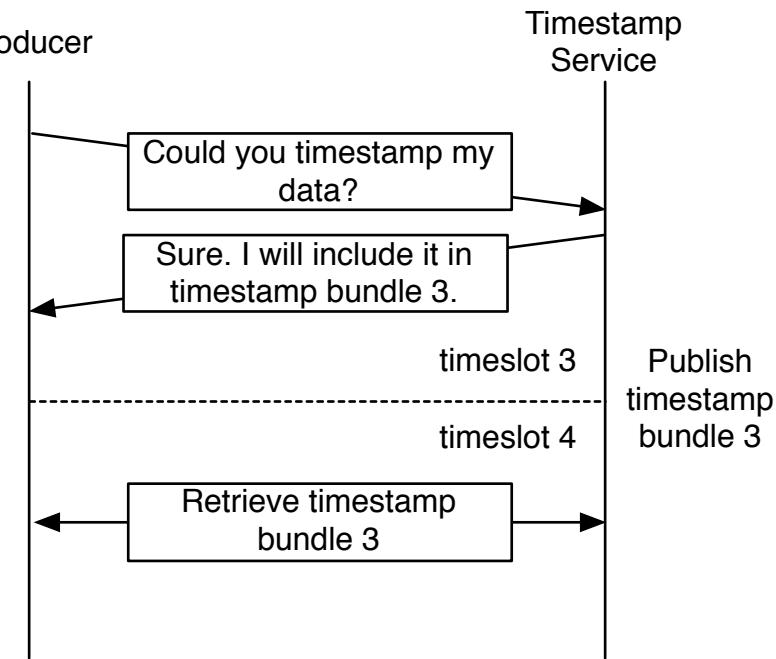
# How to rollback the clock?

- Timestamp service
  - producer requests timestamp of data from the service
  - provide existence proof of data at a given time point
- Design challenges
  - how to tell the timestamp service is honest?
  - scale with the number of timestamp records



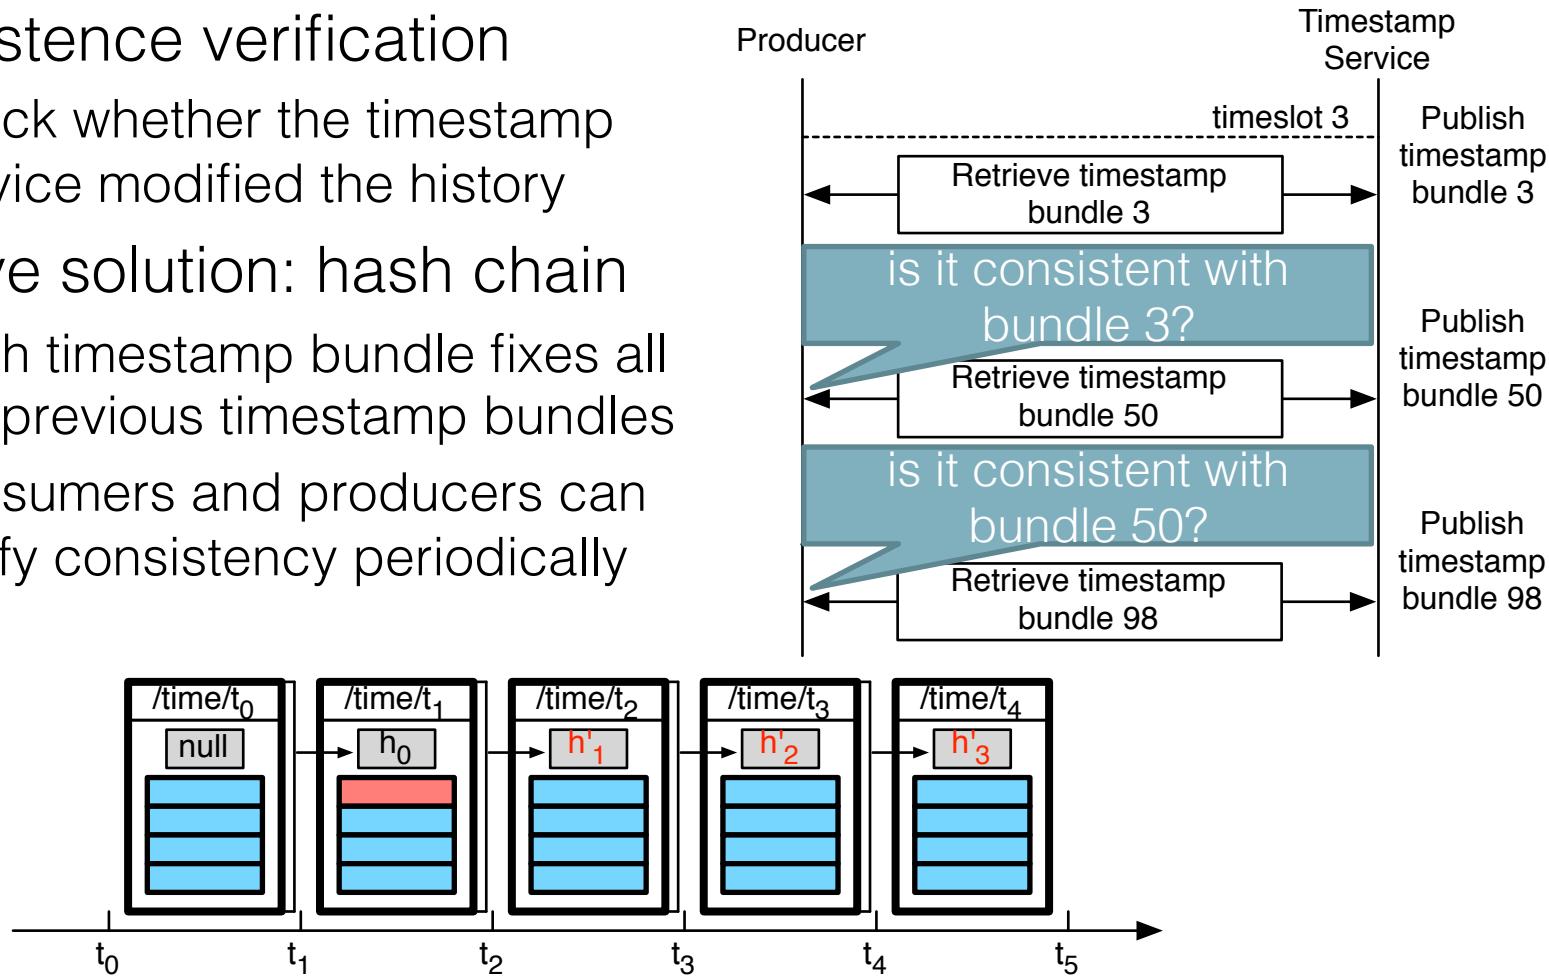
# Verifiable timestamp

- Timestamp service periodically publishes a timestamp bundle
  - containing data received during the time period
- Producer requests including its data in a bundle
- Existence verification
  - check whether data is in the corresponding bundle



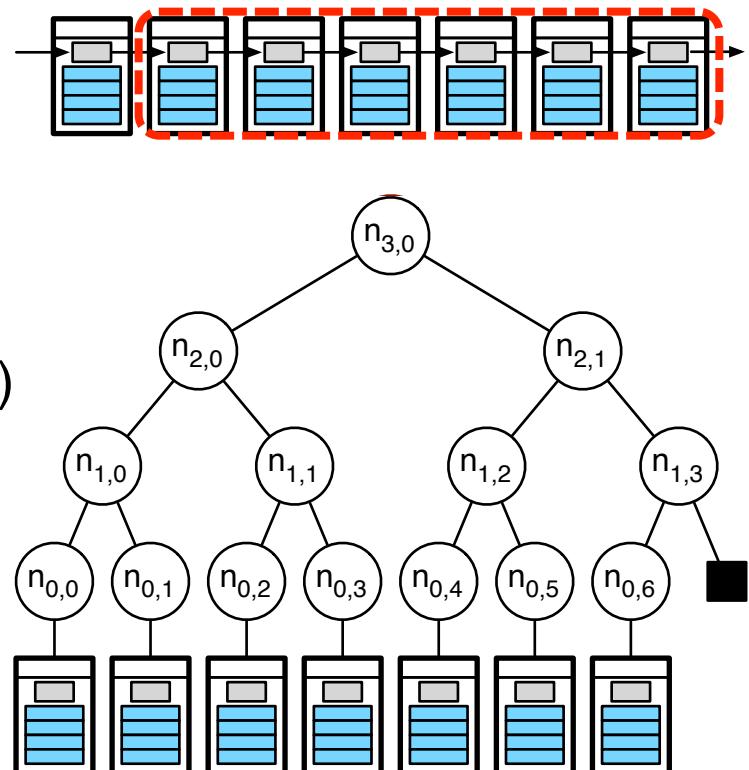
# Consistent timestamp

- Consistency verification
  - check whether the timestamp service modified the history
- A naïve solution: hash chain
  - each timestamp bundle fixes all the previous timestamp bundles
  - consumers and producers can verify consistency periodically



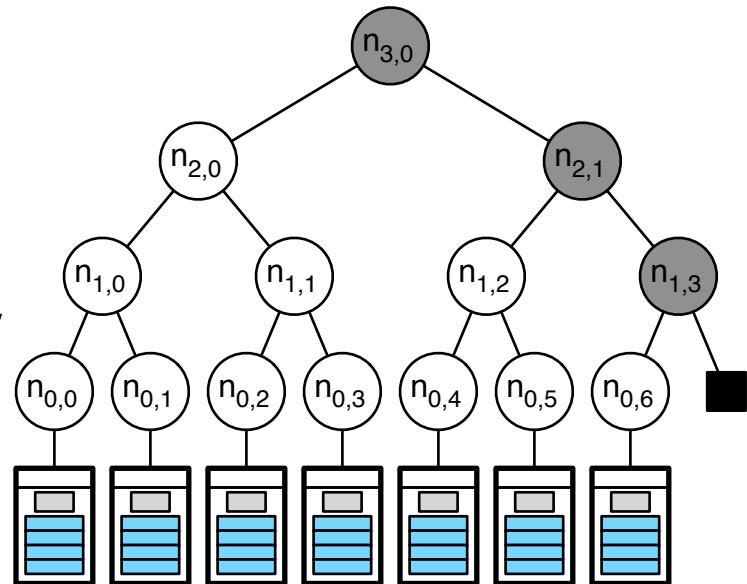
# Reduce verification overhead

- Hash chain:  $O(m)$ 
  - $m$ : number of timeslots
  - 10-min timeslots for 20 years:  $10^6$
- $k$ -ary Merkle tree:  
$$h_{i,n} = H(h_{i-1,nk^i} | h_{i-1,nk^i+1} | \dots | h_{i-1,nk^i+k-1})$$
  - root hash as the state
  - existence verification:
    - $O(\log_k m)$
  - consistency verification:
    - $O(\log_k m)$
  - 20 years timestamps
    - 4 hash computations for 32-ary Merkle tree



# Verification proof as data

- Proof is a pre-determined node set
  - simply publishes each node as data
  - consumer look up nodes necessary for verification
- Update nodes after adding a new timestamp bundle
  - complete nodes are not changed
  - at most one incomplete node at each layer

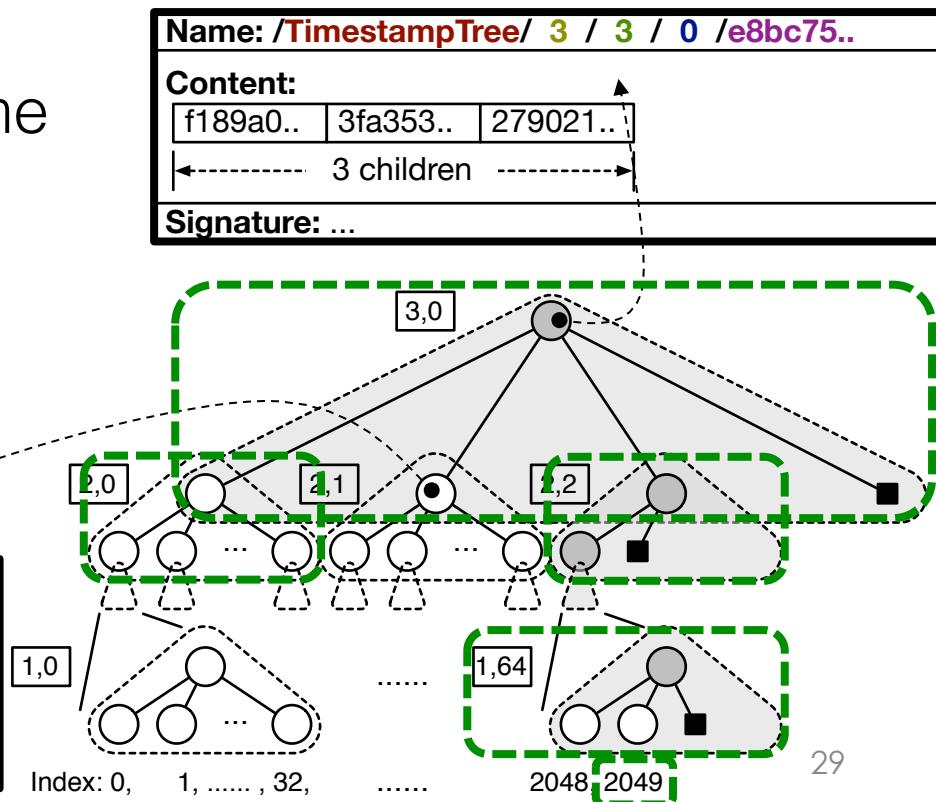


# Node data

- Naming convention
  - uniquely identify a node in a particular state  
**/[tree\_prefix]/[completeness]/[layer]/[index]/[hash]**

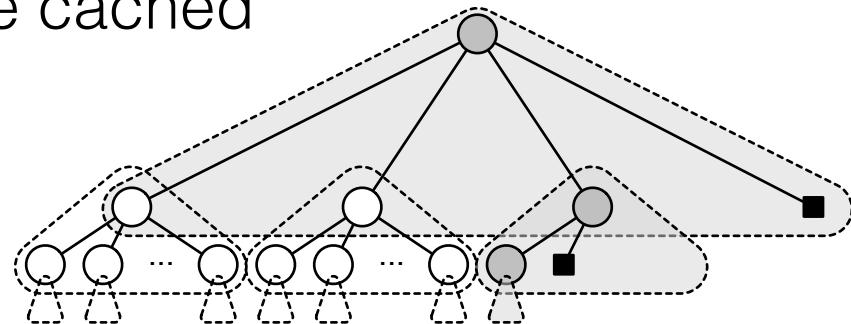
- Given a time point, the name of any node is determined  
**/TimestampTree/2050/1/64**  
**/TimestampTree/complete/2/0**  
**/TimestampTree/2050/2/2**  
**/TimestampTree/2050/3/0**

|                    |   |          |          |     |          |                    |  |  |  |
|--------------------|---|----------|----------|-----|----------|--------------------|--|--|--|
| Name:              | <b>/TimestampTree/complete/ 2 / 1 /9900a..</b>  |          |          |     |          |                    |  |  |  |
| Content:           | <table><tr><td>a2ed8b..</td><td>7ac9dd..</td><td>...</td><td>4bb231..</td></tr><tr><td colspan="4">32 children hashes</td></tr></table> | a2ed8b.. | 7ac9dd.. | ... | 4bb231.. | 32 children hashes |  |  |  |
| a2ed8b..           | 7ac9dd..  | ...      | 4bb231.. |     |          |                    |  |  |  |
| 32 children hashes |   |          |          |     |          |                    |  |  |  |
| Signature:         | ...   |          |          |     |          |                    |  |  |  |

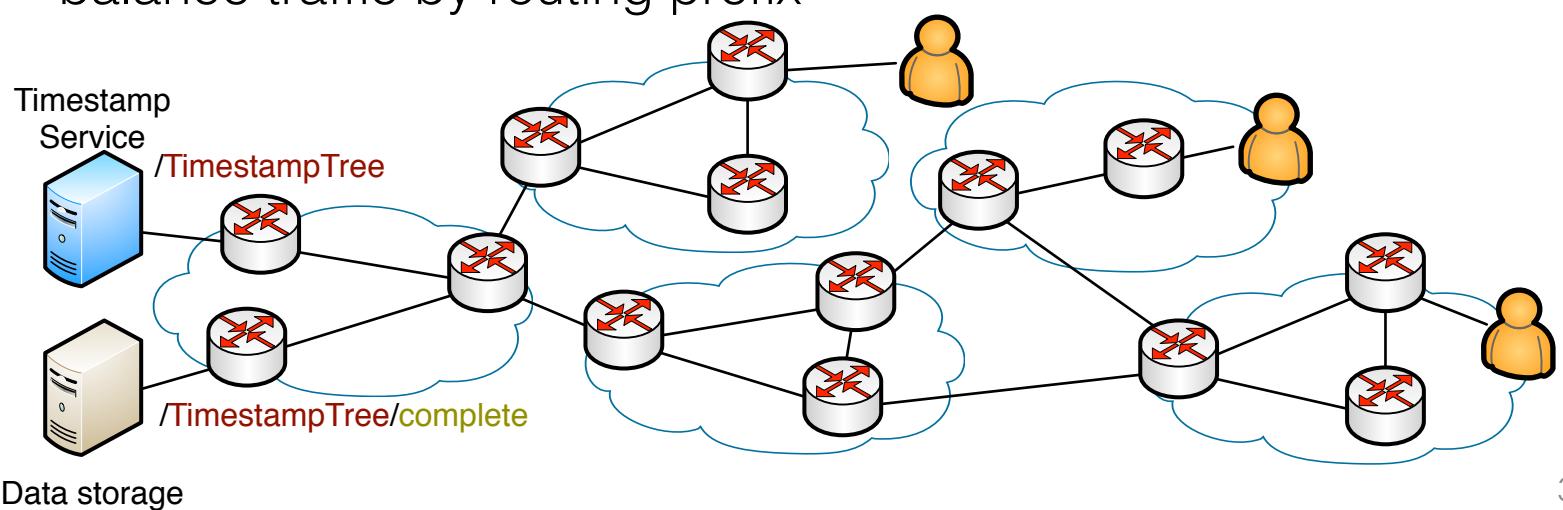


# Node retrieval

- Nodes at higher layers are cached
  - more frequently retrieved
  - root node cached almost everywhere

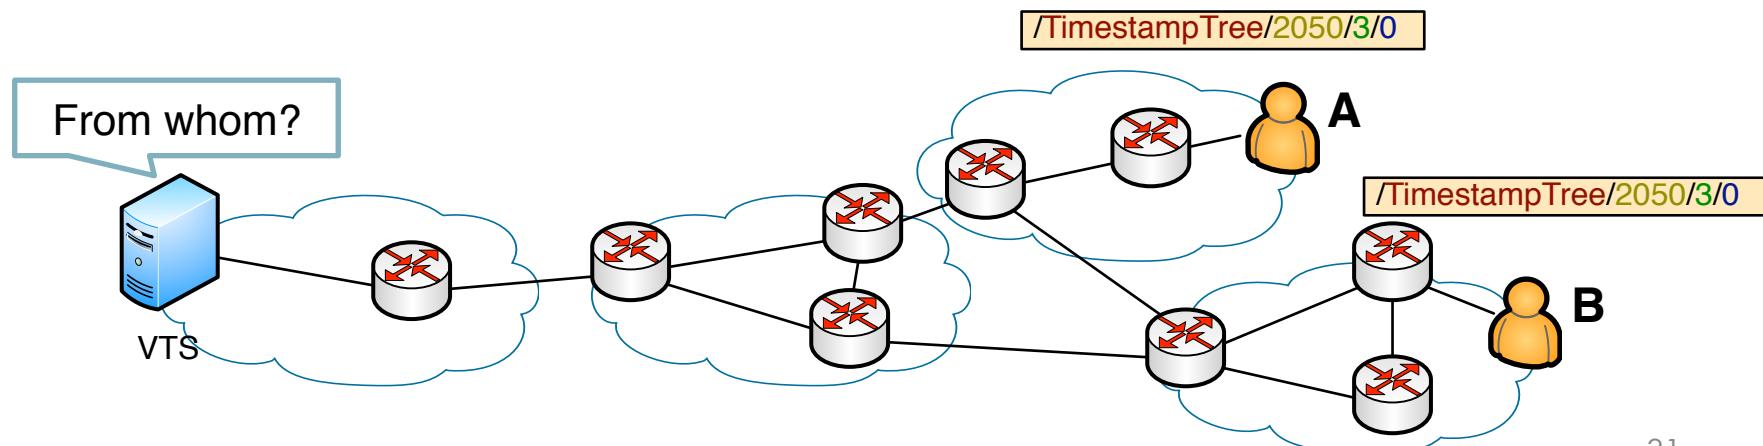


- Complete nodes can be served by dummy storage
  - balance traffic by routing prefix



# Public auditing with Merkle tree

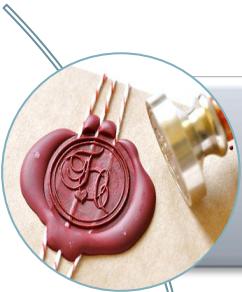
- All the users verify the consistency of timestamp service
  - occasionally retrieves the root
  - the more users, the more secure
    - single timestamp service for all the users
- Difficult to create double history
  - NDN interest does not carry sender address
  - Interest may not reach timestamp service (satisfied by cache)



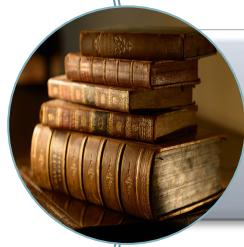
# Summary

- After fact validation is an authentication model for **non-instantaneous** communication
  - decouple the lifetime of data and signature
  - encourage the use of short-lived key
- **Untrustworthy** but **verifiable** timestamp service in NDN
  - borrow the concept public auditing concept from Certificate Transparency
  - publishing Merkle-tree as data simplifies verification query processing
  - absence of source address and efficient data distribution facilitates public auditing

# Outline



Automating Data-Centric Authenticity



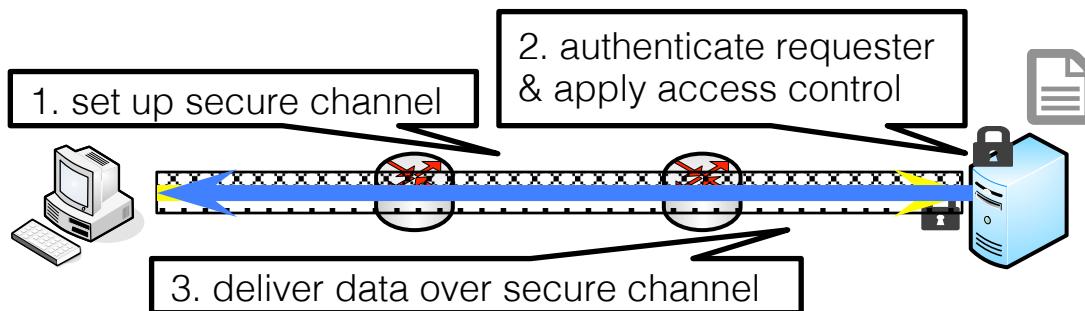
Authenticating Long-Lived Data



Automating Data-Centric Confidentiality

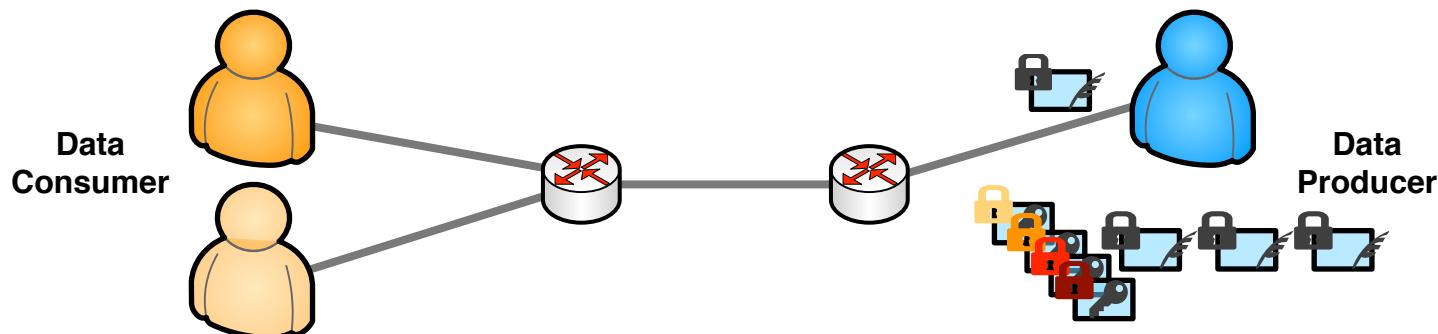
# Data confidentiality

- Current practice: perimeter-based security
  - data stays in trusted container
  - pass data to authorized users through an end-to-end secure channel
- Can we support data owner controlled confidentiality without trusted container and secured end-to-end channel?



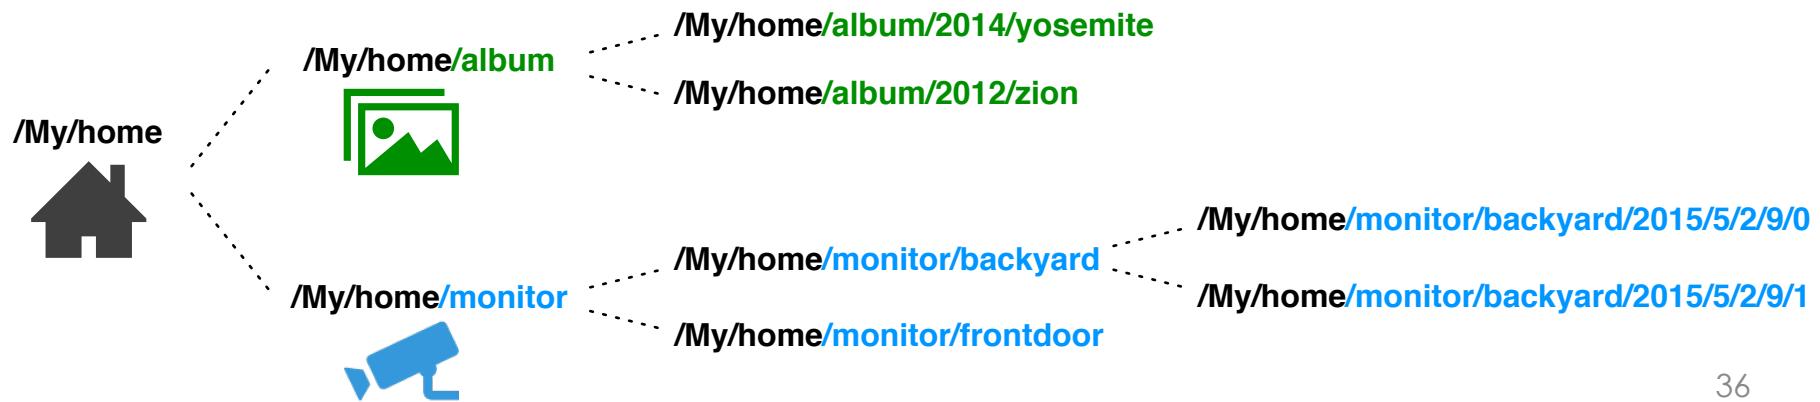
# Data-centric confidentiality

- Encrypt data at the time of production
- Distribute decryption keys to authorized consumers
- Design challenges
  - **Specify privilege using hierarchical name  
Publish encryption instruction as named keys**
  - How to distribute decryption keys efficiently?



# Distributed production & Dynamic sharing

- Shared album in SmartHome
  - members produce photos at different sites in different years
  - shared with relatives later
    - no pre-knowledge about whom the photos will be shared with
- House surveillance video
  - produced by cameras in different rooms
  - allow security personnel to watch the video when nobody at home
    - no pre-knowledge when family is out



# Content key

- Data is encrypted using a content key (C-KEY)

- symmetric key

- generated by producer

`/[content_namespace]/C-KEY`

`/My/home/album/2012/zion/C-KEY`

`/My/home/album/2012/zion/1`

`/My/home/album/2012/zion/2`

`/My/home/album/2012/zion/3`

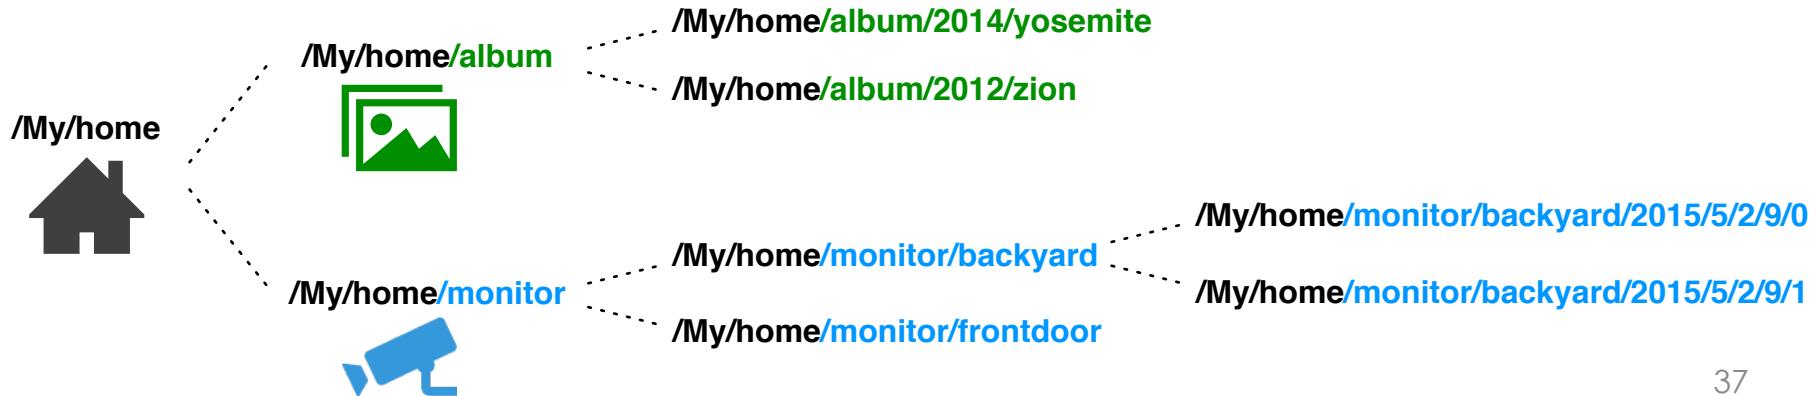


- Minimal access granularity

- encrypt data under the namespace

`/My/home/monitor/backyard/2015/5/2/9/C-KEY`

- Distributed to authorized consumers eventually

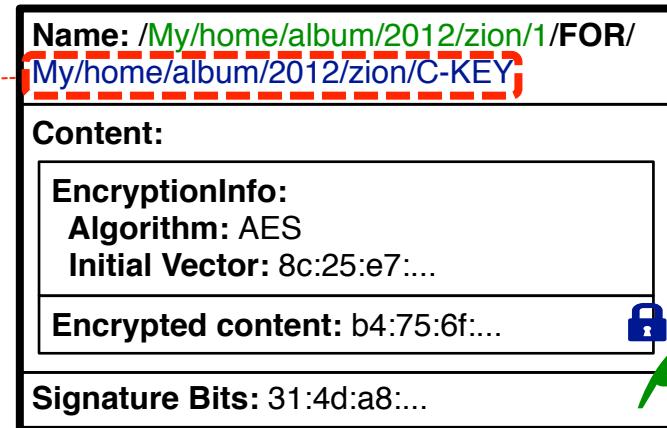


# Encrypted data

- A data packet with encrypted content
  - encryption metadata
  - encrypted content
- Encryption key name is encoded in data name  
*/[content\_name]/FOR/[encrypt\_key\_name]*
  - different keys lead to different copies of encrypted data
  - follow encryption key name, retrieve decryption key

Interest:

My/home/album/2012/zion/C-KEY/...



# Content key distribution

- Distribute content key as encrypted data
  - encrypted using authorized consumer's public key
  - producer can publish the encrypted content key later
  - consumer can construct a decryption chain following the names

Interest:

/My/home/album/2012/zion/1

Interest:

/My/home/album/2012/zion/C-KEY/FOR/My/home/relative/diane/KEY

Interest:

/My/home/album/2012/zion/C-KEY/FOR/My/home/relative/diane/KEY

Data:

/My/home/album/2012/zion/C-KEY

Data:

/My/home/album/2012/zion/C-KEY

Name: /My/home/album/2012/zion/C-KEY/  
FOR/My/home/relative/diane/KEY

Content:

EncryptionInfo:

Algorithm: AES

Initial Vector: 8c:25:e7:...

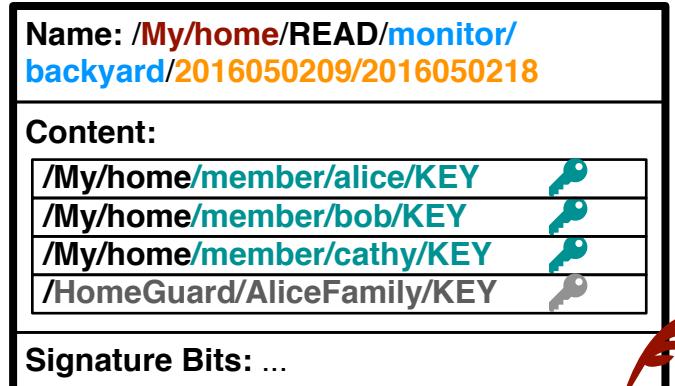
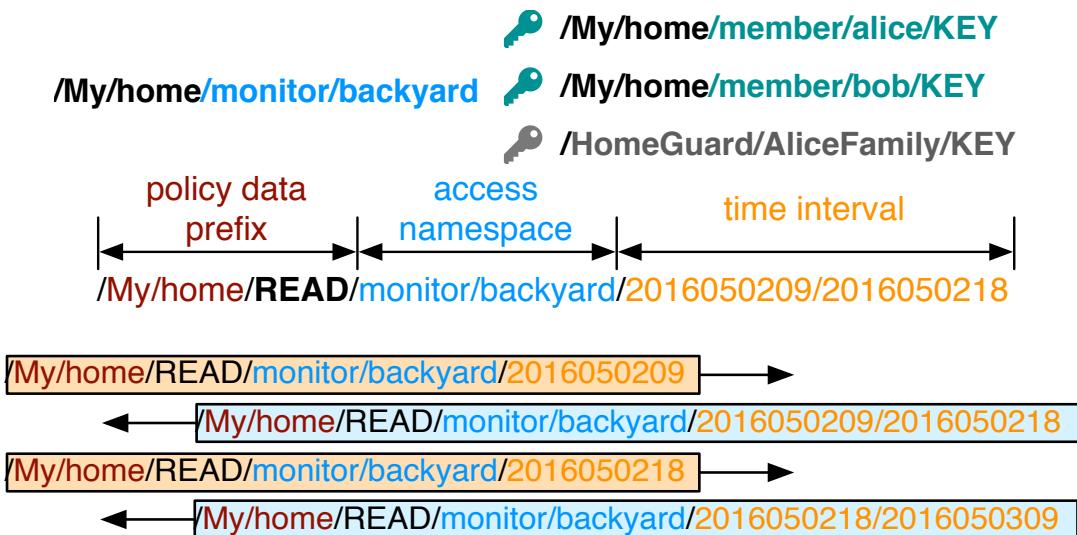
Encrypted content:



Signature Bits: 31:4d:a8:...

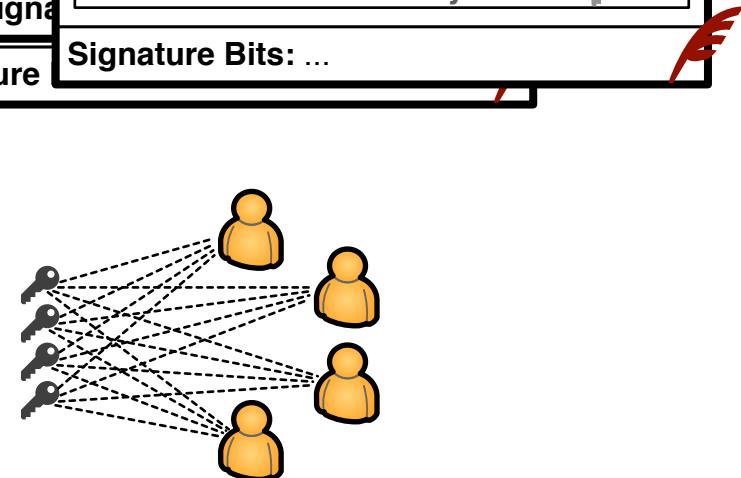
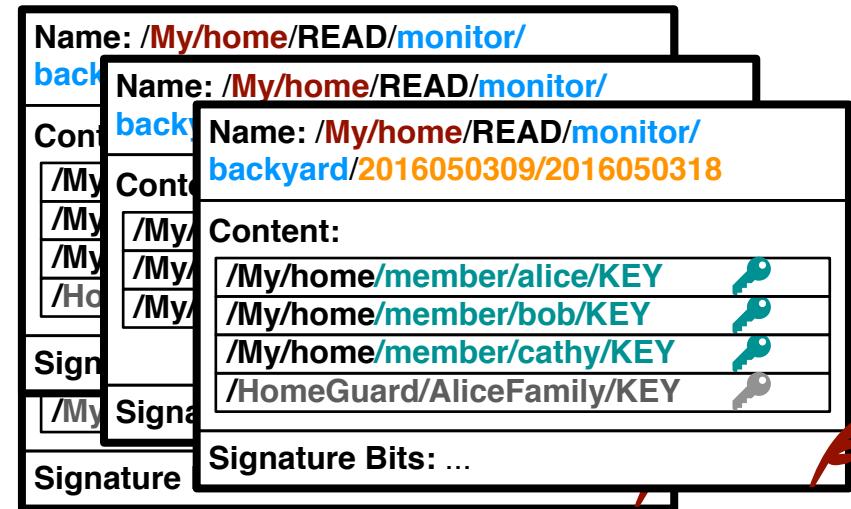
# Access control policy distribution

- Some producers require updated access control policy
  - surveillance camera
- Access control policy
  - a list of (namespace, authorized consumer key set)
- Namespace owner publishes access control policy
  - producer retrieves the latest policy



# Scalability issues

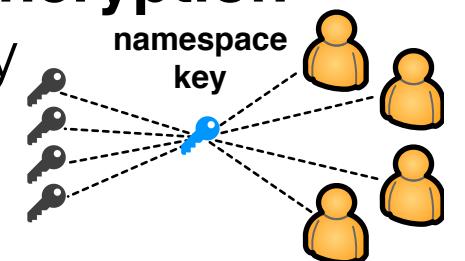
- Policy retrieval overhead
  - large data packet for popular namespace
  - redundant key retrieval
- Key encryption overhead
  - a large number of content key
  - an encrypted copy of content key for each authorized consumer
  - numbers of encrypted copies of content keys:  $O(mn)$ 
    - m: number of authorized consumers
    - n: number of content keys



**Does a producer have to know all the authorized consumers?**

# Namespace encryption key

- Namespace owner publish **namespace encryption keys** instead of namespace access policy
- Number of encrypted copies:  $O(m+n)$



Name: /My/home/READ/monitor/backyard/2016050209/2016050218

Content:

|                            |          |
|----------------------------|----------|
| /My/home/member/alice/KEY  | key icon |
| /My/home/member/bob/KEY    | key icon |
| /My/home/member/cathy/KEY  | key icon |
| /HomeGuard/AliceFamily/KEY | key icon |

Signature Bits: ...

distributed to consumers

Name: /My/home/READ/monitor/backyard/D-KEY/2016050209/2016050218/FOR/My/home/member/alice/KEY

Name: /My/home/READ/monitor/backyard/D-KEY/2016050209/2016050218/FOR/My/home/member/bob/KEY

Name: /My/home/READ/monitor/backyard/D-KEY/2016050209/2016050218/FOR/My/home/member/cathy/KEY

Name: /My/home/READ/monitor/backyard/D-KEY/2016050209/2016050218/FOR/HomeGuard/AliceFamily/KEY

Content: lock icon

Signature Bits: ...

retrieved by producers

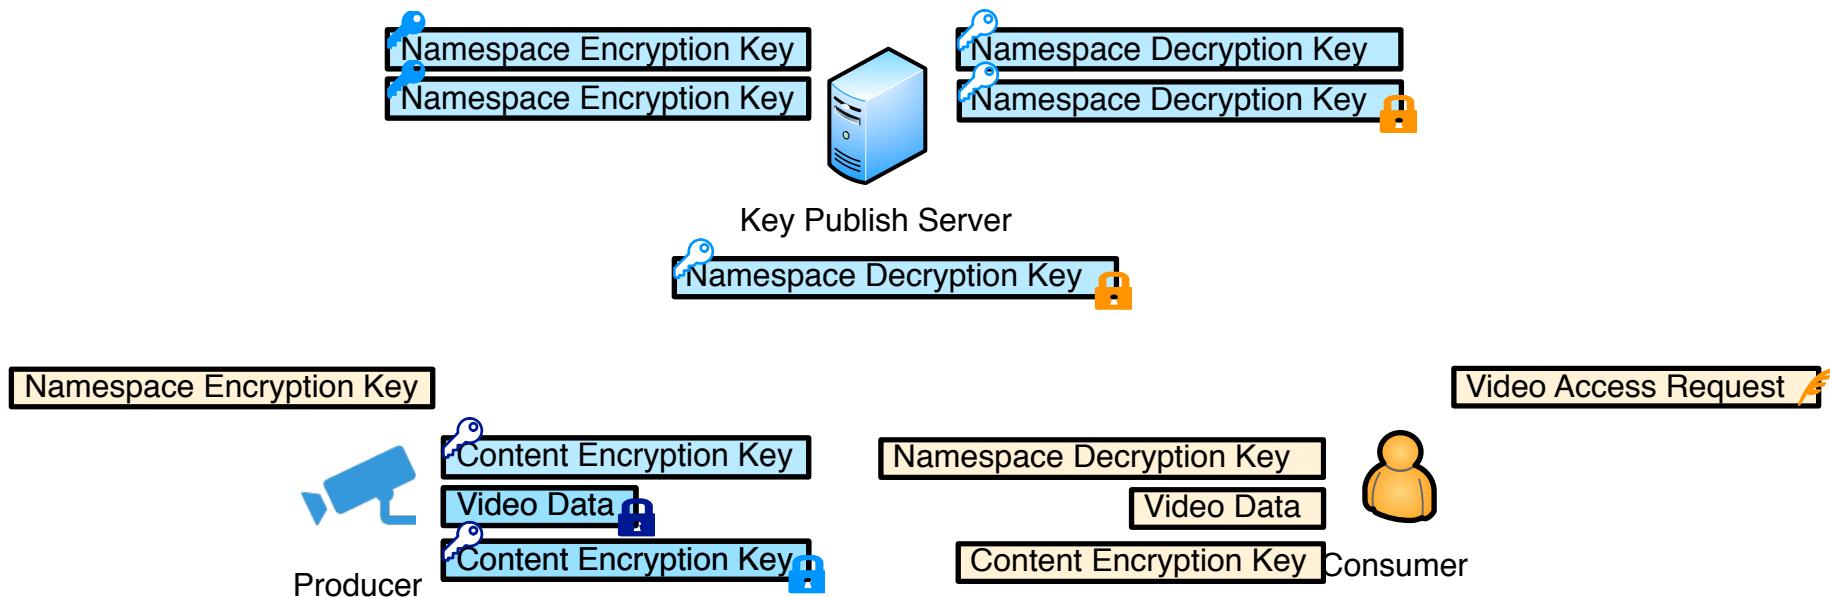
Name: /My/home/READ/monitor/backyard/E-KEY/2016050209/2016050218

Content: lock icon

Signature Bits: ...

# Automate granting access

- Namespace owner can run a key publishing server to automate data encryption
  - validate consumer's access request using trust schema
  - generate namespace decryption key for requesting consumer



# Implementation

- Available in all the NDN platform libraries
  - ndn-group-encrypt:
    - <http://github.com/named-data/ndn-group-encrypt/>
  - NDN-CCL
    - <http://named-data.net/codebase/platform/ndn-ccl/>
- Powers data access control in:
  - NDNfit: health data sharing over NDN
  - EBAMS: building management system over NDN

# Summary

- Data-centric confidentiality is a **decryption key distribution** problem
  - control access by publishing encryption/decryption keys
- Key name specifies access at **fine granularity**
  - automate data encryption
- Indirected encryption enables **scalable** key distribution

# Conclusion

- Data-centric security model enables flexible data communication model
  - reduced dependency on the data containers and channels
- Usability is critical to any security solution
  - developers need high-level abstraction
  - automation minimizes developer's workload
- Expressive names enables usable security in NDN
  - provide sufficient context and fine granularity for least privilege
  - naming pattern can represent flexible trust models and automate authentication & encryption

# Future work

- Trust schema bootstrapping
- Robust timestamp service
  - multiple instances
  - failure recovery
- Enable name confidentiality

# List of publications

- Journal and conference papers
  - Y. Yu, A. Afanasyev, D. Clark, kc claffy, V. Jacobson, and L. Zhang, "Schematizing Trust in Named Data Networking," Proc. of ACM ICN, 2015.
  - A. Afanasyev, Z. Zhu, Y. Yu, L. Wang, and L. Zhang, "The Story of ChronoShare, or How NDN Brought Distributed Secure File Sharing Back," in Proc. of IEEE MASS, 2015.
  - Y. Yu, D. Wessels, M. Larson, and L. Zhang, "Check-R: A New Method of Measuring DNSSEC Validating Resolvers," in Proc. of IEEE TMA Workshop, 2013.
  - Y. Yu, D. Wessels, M. Larson, and L. Zhang, "Authoritative Name Server Selection of DNS Caching Resolvers," in ACM Computer Communication Reviews, 2012.
- Technical reports
  - W. Shang, Y. Yu, R. Droms, and L. Zhang, "Challenges in IoT Networking via TCP/IP Architecture," Technical Report NDN-0038, 2016.
  - V. Lehman, A. Hoque, Y. Yu, L. Wang, B. Zhang, and L. Zhang, "A Secure Link State Routing Protocol for NDN", Technical Report NDN-0037, 2016.
  - W. Shang, Y. Yu, T. Liang, B. Zhang, and L. Zhang "NDN-ACE: Access Control for Constrained Environments over Named Data Networking", NDN, Technical Report NDN-0036, 2015
  - Y. Yu, A. Afanasyev, and L. Zhang "Name-Based Access Control", Technical Report NDN-0034, 2015
  - Y. Yu "Public Key Management in Named Data Networking", Technical Report NDN-0029, 2015
  - Y. Yu, A. Afanasyev, Z. Zhu, and L. Zhang "An Endorsement-based Key Management System for Decentralized NDN Chat Application", Technical Report NDN- 0023, 2014
  - Y. Yu, J. Cai, E. Osterweil, and L. Zhang "Measuring the Placement of DNS Servers in Top-Level-Domain" Technical Report, May. 2011

# Special Thanks To Collaborators



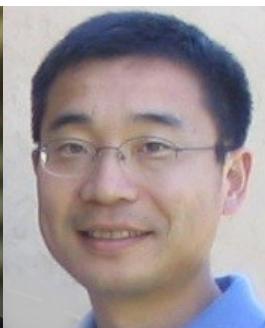
Lixia Zhang



Van  
Jacobson



Alex  
Halderman



Beichuan  
Zhang



Lan Wang



David Clark



kc claffy



Alexander  
Afanasyev



Zhenkai Zhu



Wentao  
Shang



Haitao Zhang



Spyridon  
Mastorakis



Qiuhan Ding

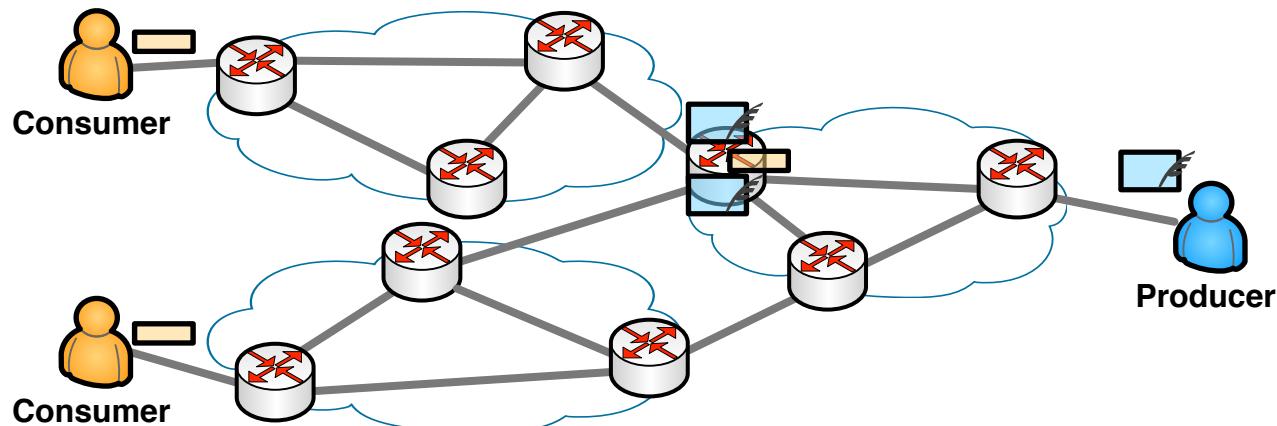


Prashanth  
Swami

# **End**

# NDN Overview

- Native multicast
  - Interest for the same data can be merged

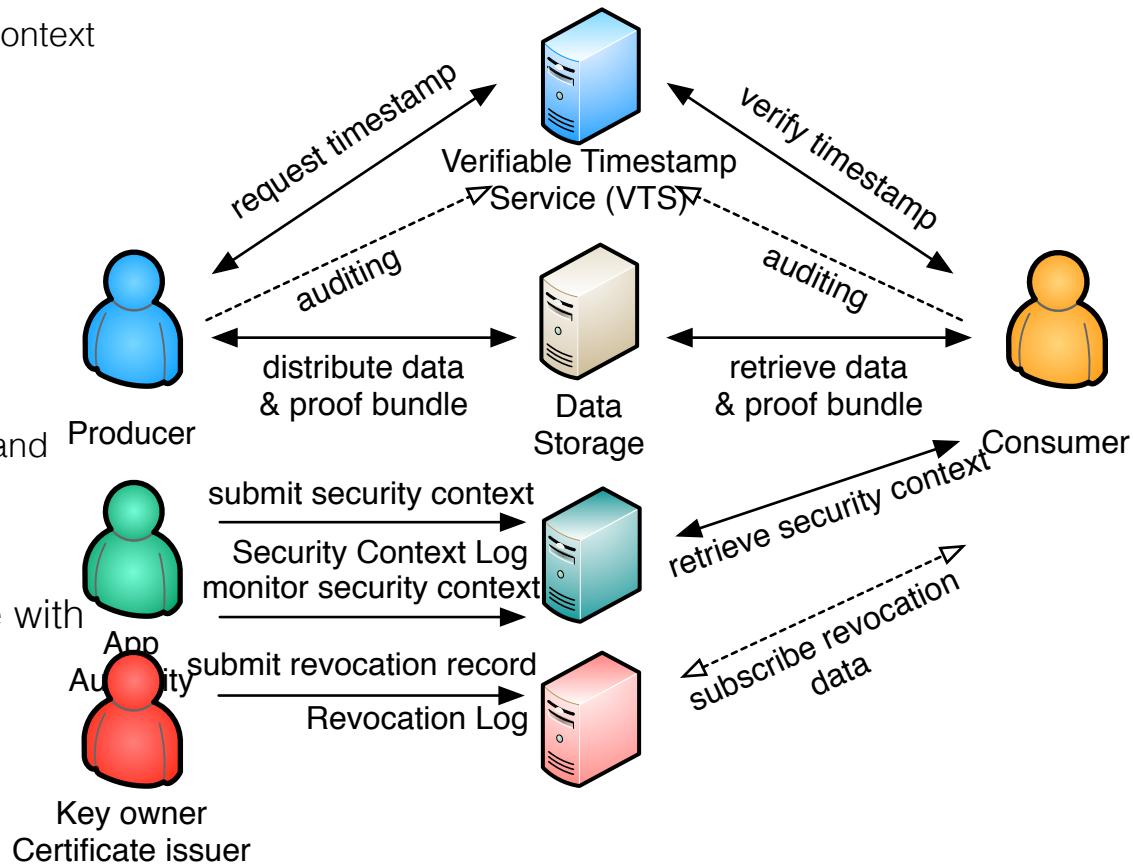


# Data-centric security & NDN

- Named Data Networking (NDN)
  - data-centric communication primitives
  - retrieve data by name rather by host
- NDN enables data-centric security
  - per-packet signature
  - hierarchical naming
    - security context
    - least privilege
  - efficient key distribution

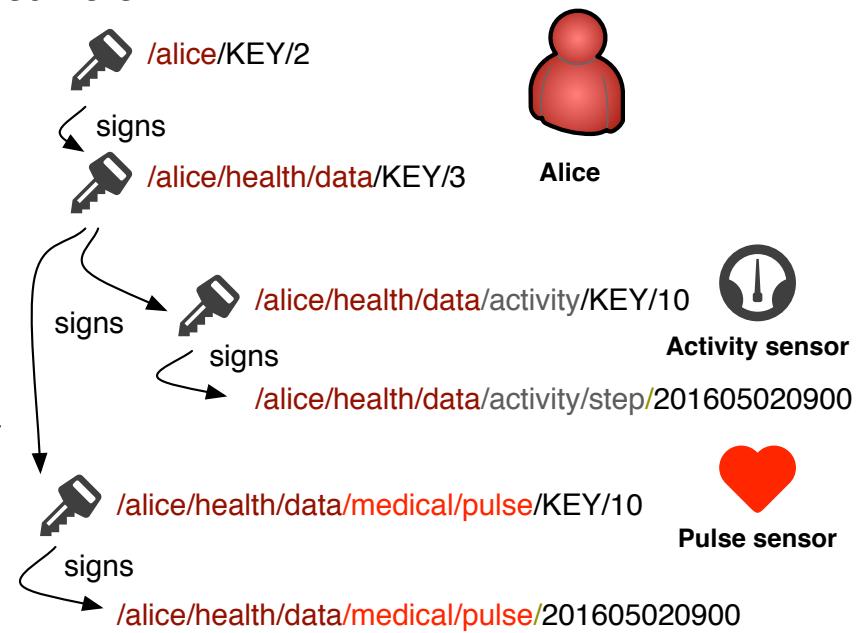
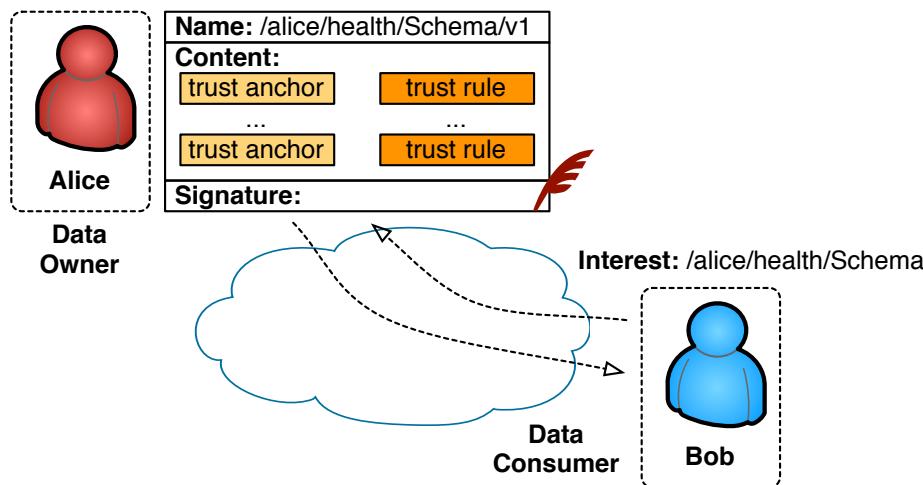
# SigLogger Overview

- Security Context Log
  - record security context over the time
    - trust schema
  - assure only one valid version of context at any time point
  - secure through publicity
- Revocation Log
  - record revocation over the time
  - promptly distribute revocation information to consumers
- Verifiable Timestamp Service
  - provide existence proof of data (and keys)
  - untrustworthy but auditable
- Producer distributes proof bundle with data
  - timestamp of data
  - intermediate keys
  - timestamp of keys



# Signing-based write access

- Key name represents capability
  - capable of producing data under a namespace
  - capable of delegating the write access of a sub-namespace to others
  - signing key hierarchy
- Express write access control policy as trust schema
- Distribute trust schema as data
  - published by data owner retrieved by consumers



# Append-only timestamp service

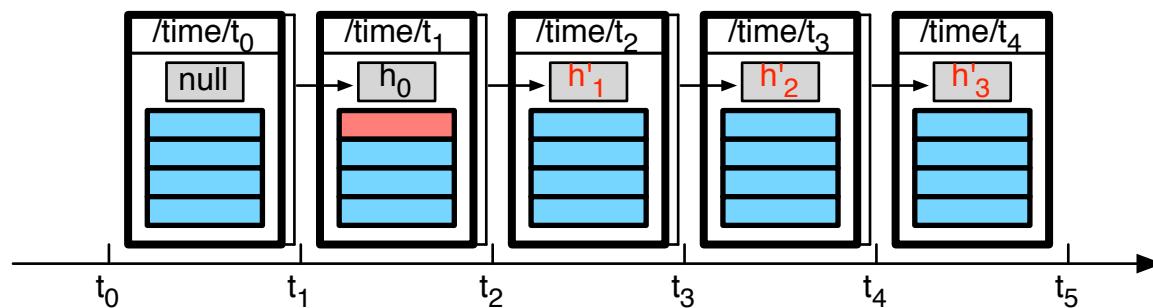
- Chaining timestamp data by hash
  - each timestamp data fixes all the previous timestamp bundle

$$h_k = H(TB_k || h_{k-1})$$

- Consistency verification

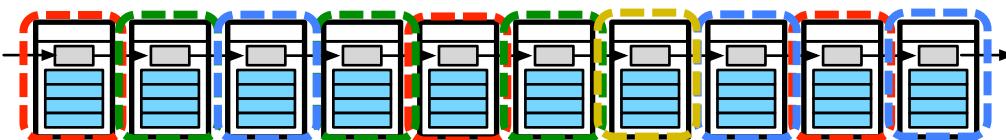
• **No way to know whether the timestamp service is honest about  $[t_i, t_j]$**

- any modification before  $t_i$  will be detected



# Public auditing

- Easy to catch misbehavior if
  - the consistency of each timestamp bundle is checked by at least one consumer
- Each consumer verifies consistency occasionally
- A lot of consumers collectively audit the single timestamp service
- How to minimize verification overhead



# Use Case Example

- Alice obtains a timestamp for her thesis
  - also for keys if not timestamped
  - distribute data with keys and timestamps
- Bob verifies the existence of keys using timestamp
  - verifies data using keys

