

Support Mobile and Distributed Applications with Named Data Networking

Zhenkai Zhu

Computer Science Department
University of California, Los Angeles
CA, 90095

zhenkai@cs.ucla.edu

May 22, 2013

Supporting mobility in the global Internet

- thoroughly studied existing IP mobility solutions
- developed solution for the scalability problem in Global HA to HA protocol
- provided a new perspective of mobility support

Exploring new design patterns for distributed applications over NDN

- explored naming conventions in application design
- proposed a general-purpose dataset synchronization protocol that has the potential of supporting a wide range of distributed applications
- developed security solutions for distributed applications using NDN's data-centric approach

- Today's Internet is increasingly mobile
- IP's host-to-host communication model is increasing challenged by emerging communication patterns
- Despite years of IP mobility support research, few solutions enjoyed wide adoption
 - it is time to rethink what is mobility support and how to align mobility support with the applications' need
- NDN aims to accommodate emerging communication patterns by taking fundamental changes to today's IP Internet architecture
 - we must learn how to design applications to fully exploit the benefits of such an architectural shift by trying out new applications
- This work is to push forward this new front of networking research

Supporting mobility in the global Internet

The IP mobility problems

- How to find a mobile after it moves?
- How to minimize the interruption to the communications?
- How to make sure one is talking to the right host?

- One can reach a mobile by the same IP address regardless of whether it moves or not
 - e.g. broadcast the mobile's location via routing
- One can find the mobile's location based on a stable piece of information
 - e.g. home address, DNS name
 - need to maintain an up-to-date mapping

Three basic components of IP mobility support


- Stable identifier
 - sender's knowledge about the mobile that does not change due to mobility
- Locator
 - an IP address for the mobile's current location
- Mapping between the two
 - kept in dedicated network entities
 - or by the routing system

IP mobility solutions


I conducted a complete and systematic survey of IP mobility solutions that have been proposed in order to gain a good understanding of the solution space and shed light on future efforts.

Protocol	Year	Protocol	Year
Columbia	1991	TIMIP	2001
Virtual IP	1991	M-SCTP	2002
LSR	1993	HIP	2003
Mobile IP	1996	Connexion	2004
MSM-IP	1997	ILNPv6	2005
Cellular IP	1998	Global HAHA	2006
HMIP	1998	PMIP	2006
FMIP	1998	BTMM	2007
HAWAII	1999	WINMO	2008
NEMO	2000	LISP-Mobility	2009
E2E	2000		


Mapping




Routing



Mapping



Routing



Summary of IP mobility solution space

- Mobility support essentially involves three basic components
 - a stable identifier, a locator and the mapping between the two
- Broadcast (routing) based approach is simplest, most robust, and work well in small networks, but raises scaling concerns when applied to large networks
- Mapping based approach (home agent, DNS, etc.) can scale well with large network, but introduces dependency on a 3rd party.
- Managing local movements local helps improve the performance and scalability

Limitations of IP mobility support


- Security is tied to IP addresses
 - often weak and need frequent updates during the move
 - quite a few proposals didn't even mention about security
- Incomplete solution: only considers cases when mobiles are connected to the infrastructure
 - yet mobility can easily lead to intermittent connectivity or ad hoc connectivity among a set of mobiles
 - ad hoc mobility and DTN becomes separate branches in networking research with their own solutions
- Because IP is for point-to-point delivery, so far mobility solutions also focus only on maintaining point-to-point sessions

A new perspective on mobility support

What NDN brings


- Data-centric security
 - every piece of data is named and signed
 - where the data is obtained doesn't matter
- Receiver-driven communications
 - must send Interest for data
 - data follows the reverse path back to requester
- Intelligent data plane
 - per-packet state at routers

Interest	Data
Name	Name
Selectors	Content
Nonce	Signed Info
	Signature




Security done right

- Security has always been a big challenge for IP mobility solutions
- NDN secures data
 - decouples trust in data from trust in host




(a) IP




(b) NDN

Enhanced performance with network caching


- Data can be easily cached anywhere
 - 1 fast recovery for packet drops at last (wireless) hop
 - 2 smooth handoff without special optimization
 - 3 always fetch popular content from nearest cache or repository



(1) last hop recovery




(2.a) optimization for handoff in IP



(2.b) handoff in NDN


Naturally support DTN and ad hoc networks

- NDN already provides two basic building blocks for DTN and ad hoc networks
 - a way to identify data → every piece of data is named
 - a way to carry data around → data is secured and can be cached anywhere
- Furthermore, NDN fully utilizes the broadcast nature of wireless communications
 - Interest can be broadcasted and whoever has the data can reply
 - Overheard data can be cached




- IP is about connectivity
 - has to acquire IP address and report IP address changes so that the other end can “push” data

- In NDN, data flows back along the reverse path of the Interest
 - no need to acquire names
 - no need to report location changes




- The network needs to route Interests based on data names
 - producers whose names are not on FIB must have a way for Interests to reach them
- The lessons we learned from IP mobility research can be applied
 - stable identifier: the name space assigned to a producer
 - e.g. /ndn/ucla.edu/cs/foo
 - locator: the name prefix that hints the location of a producer
 - e.g. /ndn/stanford.edu
 - mapping:
 - broadcast: simplest and robust, and takes advantage of broadcast nature of wireless communications
 - DNS: one can do DNS lookup to find the mobile

An example of DNS based support for mobile producer




Flexible data fetching

- NDN's *forwarding strategy* provides the much desired flexibility in fetching mobile producer's data
 - enabled by packet statistics and loop-free Interest forwarding
 - also take into considerations of each face's properties (e.g. broadcast interface, traffic cost) and namespace



(a) utilizing multiple interfaces



(b) exploring alternative path

Publications on mobility support

- Supporting mobility in the global Internet (ACM MobiCom MICNET Workshop '09)
- Supporting mobility for Internet cars (IEEE Communications Magazine 2011)
- SAIL: A scalable approach for wide-area IP mobility (IEEE Infocom MobiWorld Workshop '11)
- Home as you go: an engineering approach to mobility-capable extended home networks (AINTEC '11)
- A survey of mobility support in the Internet (IETF RFC 6301)
- Understanding Apple's back to my mac service (IETF RFC 6281)
- A new perspective on mobility support under submission to MobiArch '13

Exploring new design patterns for distributed applications over NDN


Why new design patterns

- NDN uses application defined names for data delivery
 - so apps can now use names to achieve new functions that they were unable to do before
- Exactly how to design apps to take such advantages is yet to be explored
- NDN research has been taking on an app-driven approach to get the answer

Distributed applications in today's Internet


■ Centralized designs

- centralized control
- single point of failure
- inefficient data distribution




■ Peer-to-peer designs

- efficient data distribution may not be the goal
- mis-matching between P2P overlay and underlying topology




Distributed applications over NDN: Goals


- Robust, fully decentralized
- Efficient data distribution
- Mobile friendly
- Secure



A simple distributed application: group text chat



A simple distributed application: group text chat



The problem is really synchronizing the chat dataset among the three.

- The need is ubiquitous in distributed applications
 - retrieving all messages to a chatroom
 - collecting a list of conferences and speakers in audio conferencing
 - collecting all revisions in collaborative editing
- The problems:
 - synchronizing a dataset produced by multiple parties in an unpredictable pattern
 - no pre-knowledge about parties

The role of naming

- Naming is an important aspect of NDN application design
- Names determine how Interests are forwarded in the network
 - and also identifies the application process if it reaches the producer
- Proper naming can often simplify application designs

Naming in dataset synchronization

- Use broadcast namespace for rendezvous
 - Interest carrying broadcast name will be broadcasted in a domain
 - reaches all potential data producers

/ndn/broadcast/chronos/lunch-talk/4b01...



- Use routable namespace for actual data items
 - use broadcast names for all data would be prohibitively expensive


/wonderland/alice/chronos/lunch-talk/791



- Simplify the knowledge about a producer's data
 - e.g., name the data of a producer sequentially


ChronoSync: decentralized dataset state synchronization

- Maintain knowledge of a dataset by tracking what each producer has produced so far




ChronoSync: decentralized dataset state synchronization

- Maintain knowledge of a dataset by tracking what each producer has produced so far
- Compactly represent a user's knowledge of the dataset in a hash, or state digest



ChronoSync: decentralized dataset state synchronization


- Maintain knowledge of a dataset by tracking what each producer has produced so far
- Compactly represent a user's knowledge of the dataset in a hash, or state digest
- Exchange state digests among all users using broadcast (sync) Interests



/ndn/broadcast/chronos/lunch-talk/4b01...

— (1) —|— (2) —|— (3) —|—

Tracking dataset state: digest-tree




- Each child node maintains the knowledge of a producer
 - sorted according to the name prefix of each producer
- Hash children nodes in order to produce state digest
- When a producer's state changes, update the branch and re-calculate state digest

Tracking dataset state: digest-log

State Digest	Changes
0000...	Null
9w35...	[Alice's prefix, 1]
...	...
23ab...	[Bob's prefix, 31], [Alice's prefix, 19]
05t1...	[Bob's prefix, 32]


- Maintaining the history of state digest changes
 - useful when old state digest is received, e.g. from a user who just recovers from a temp disconnection
- Not essential for the correctness of ChronoSync
 - application can set an upper bound of log size

Stable state




- Everyone has the same knowledge of a dataset
- Sync Interests carry an identical state digest
 - regardless of number of users in a group, at most one sync Interest sent over a link in one direction

Disseminate knowledge of new data




- After producing new data, a producer observes different state digest
 - reply the pending sync Interest with the new data info
- Others update their state digest once receiving the reply
 - producing the same new state digest

Disseminate knowledge of new data




- After producing new data, a producer observes different state digest
 - reply the pending sync Interest with the new data info
- Others update their state digest once receiving the reply
 - producing the same new state digest

Handle simultaneous data generation



- Problem: one sync Interest can only retrieve one sync data
 - users are partitioned into groups with different knowledge about the dataset
- Optimization: use exclude filter
 - need to exclude at most $(N - 1)$ known sync data if N producers generated data simultaneously

Problem: network partitions



State Digest	Changes
...	...
9w35...	[Alice's prefix, 1]
23ab...	[Bob's prefix, 31]
07t1...	[Alice's prefix, 2]
ks23...	[Alice's prefix, 3]

Alice and Bob's view

State Digest	Changes
...	...
9w35...	[Alice's prefix, 1]
ux53...	[Ted's prefix, 9]
ik29...	[John's prefix, 0]
990s...	[John's prefix, 2]

John and Ted's view


Handling network partitions

- Any set reconciliation algorithm can be used
- A simple recovery method is used as default
 - explicitly asking for information from the users who produced the unknown state digest
 - recovery data contains the whole snapshot of the dataset state (i.e. information about each data producer)

/ndn/broadcast/chronos/lunch-talk/recovery/ks23...




Overview of ChronoSync



ChronoShare: distributed file sharing with ChronoSync

- Starting from an empty shared folder, the state of the folder is altered whenever an action from a participant is applied
- The state of the folder can be determined by applying all actions from all participants in a deterministic way
- ChronoSync can be used track the actions




Evaluation

- Goals: to confirm

- ChronoSync synchronize dataset state quickly and efficiently
- ChronoSync is robust again network partitions, link failures and packet losses


- Methodology

- simulation in ndnSIM module of NS-3
- evaluate performance of ChronoChat, a group text chat app implemented based on ChronoSync
- use central server based TCP/IP implementation of IRC as baseline



Evaluation setup


- Topology: 52 nodes, 84 links [1]
 - 100 Mbps bandwidth for each link
- Traffic: 1000 messages following exponential distribution with mean of 5 seconds [2]
- All nodes participate in one chatroom




- [1] N. Spring et al., Measuring ISP topologies with Rocketfuel.
- [2] C. Dewes et al., An analysis of Internet chat systems.

- Synchronization delay: the time needed for the last user to learn a new data piece after its generation
- Communication pairs: the pairs of users that are still able to chat in face of network failures
- Traffic pattern: how is the traffic distributed over the links
- Total overhead: how many packets are transmitted


Synchronization delay




Synchronization delay under lossy environment




Resilience to network failures



Traffic pattern




Total overhead



Discussion: broadcast in large networks

- Broadcast directly in large networks is costly
- A broadcast overlay can dramatically reduce the cost



- ACT: an audio conference tool over Named Data Networking ACM SIGCOMM ICN Workshop '11
- A new approach of securing audio conference tools AINTEC '11
- Chronos: serverless multi-user chat over NDN NDN Tech-Report 0008, 2012
- Let's ChronoSync: decentralized dataset state synchronization in Named Data Networking under submission

- ACT: multiple party audio conference tool
- XMPP-MUC: multi-user text chat
- ChronoChat: Chronosync based multi-user text chat
- ChronoShare: ChronoSync based file sharing application
- peets: WebRTC based audio/video conferencing
- ChronoSync: C++ library
- py-chronos: python API for ChronoSync
- NDN.cxx: C++ API for NDN
- ndndump: packet analyzer for NDN

Concluding thoughts

- The Internet needs changes to accommodate billions of mobile devices and dramatically changed communication patterns.
- By aligning mobile communications with application's data-centric nature, we provided a completely different perspective on mobility support
- ChronoSync leverages NDN's naming and data multicast capability to achieve efficient and robust dataset state synchronization
 - effectively names the state of a dataset by its given digest at a given time
 - removes single point of failure and traffic concentration problems
 - provides useful building blocks in supporting distributed applications
- Look forward, we hope this work can help stimulate more discussions on the design space of mobile and distributed applications over NDN

Acknowledge

- Alexander Afanasyev made significant contribution to the design and implementation of ChronoSync. He is also a collaborator on ChronoShare and NDN.cxx.
- Chaoyi Bian made significant contribution to XMPP-MUC